

```
from google.colab import files
uploaded = files.upload() # This will open a file picker
```

Choose Files | cleaned\_features.csv  
**cleaned\_features.csv**(text/csv) - 545825 bytes, last modified: 20/11/2025 - 100% done  
Saving cleaned\_features.csv to cleaned\_features.csv

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import io

# Make plots look nice
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")
%matplotlib inline
```

```
# Load the CSV file you just uploaded
df = pd.read_csv('cleaned_features.csv')

# Show first 5 rows
df.head()
```

	Store	Date	Temperature	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkDown5	CPI	Unemployment	IsH
0	1	2010-02-05		42.31	2.572	NaN	NaN	NaN	NaN	NaN	211.096358	8.106
1	1	2010-02-12		38.51	2.548	NaN	NaN	NaN	NaN	NaN	211.242170	8.106
2	1	2010-02-19		39.93	2.514	NaN	NaN	NaN	NaN	NaN	211.289143	8.106

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# See shape (rows, columns)
print("Shape:", df.shape)
```

Shape: (8190, 12)

```
# See column names and data types
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8190 entries, 0 to 8189
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Store        8190 non-null   int64  
 1   Date         8190 non-null   object 
 2   Temperature  8190 non-null   float64
 3   Fuel_Price   8190 non-null   float64
 4   MarkDown1    4032 non-null   float64
 5   MarkDown2    2921 non-null   float64
 6   MarkDown3    3613 non-null   float64
 7   MarkDown4    3464 non-null   float64
 8   MarkDown5    4050 non-null   float64
 9   CPI          7605 non-null   float64
 10  Unemployment 7605 non-null   float64
 11  IsHoliday    8190 non-null   bool    
dtypes: bool(1), float64(9), int64(1), object(1)
memory usage: 712.0+ KB
None
```

```
# See basic statistics
df.describe()
```

	Store	Temperature	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkDown5	
<b>count</b>	8190.000000	8190.000000	8190.000000	4032.000000	2921.000000	3613.000000	3464.000000	4050.000000	7605.000000
<b>mean</b>	23.000000	59.356198	3.405992	7032.371786	3384.176594	1760.100180	3292.935886	4132.216422	172.461000
<b>std</b>	12.987966	18.678607	0.431337	9262.747448	8793.583016	11276.462208	6792.329861	13086.690278	39.731000
<b>min</b>	1.000000	-7.290000	2.472000	-2781.450000	-265.760000	-179.260000	0.220000	-185.170000	126.061000
<b>25%</b>	12.000000	45.902500	3.041000	1577.532500	68.880000	6.600000	304.687500	1440.827500	132.361000
<b>50%</b>	23.000000	60.710000	3.513000	4743.580000	364.570000	36.260000	1176.425000	2727.135000	182.761000
<b>75%</b>	34.000000	73.880000	3.743000	8923.310000	2153.350000	163.150000	3310.007500	4832.555000	213.931000
<b>max</b>	45.000000	101.950000	4.468000	103184.980000	104519.540000	149483.310000	67474.850000	771448.100000	228.971000

```
# Convert Date column to proper datetime
df['Date'] = pd.to_datetime(df['Date'])

# Convert IsHoliday to 1/0 (True → 1, False → 0)
df['IsHoliday'] = df['IsHoliday'].astype(int)

# Sort by date (important!)
df = df.sort_values('Date').reset_index(drop=True)

# Check again
df.head()
```

	Store	Date	Temperature	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkDown5	CPI	Unemployment	IsHoliday
0	1	2010-02-05	42.31	2.572	NaN	NaN	NaN	NaN	NaN	211.096358	8.106	0
1	16	2010-02-05	19.79	2.580	NaN	NaN	NaN	NaN	NaN	189.381697	7.039	0
2	31	2010-02-05	39.05	2.572	NaN	NaN	NaN	NaN	NaN	210.752605	8.324	0

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# Fill missing MarkDown values with 0 (meaning no promotion that week)
markdown_cols = ['MarkDown1', 'MarkDown2', 'MarkDown3', 'MarkDown4', 'MarkDown5']
df[markdown_cols] = df[markdown_cols].fillna(0)

# Check how many missing values left
print(df.isnull().sum())
```

```
Store          0
Date           0
Temperature    0
Fuel_Price     0
MarkDown1      0
MarkDown2      0
MarkDown3      0
MarkDown4      0
MarkDown5      0
CPI            585
Unemployment   585
IsHoliday       0
dtype: int64
```

```
from google.colab import files
uploaded = files.upload() # This will open a file picker
```

Choose Files | train.csv  
**train.csv**(text/csv) - 12842546 bytes, last modified: 20/11/2025 - 100% done  
 Saving train.csv to train.csv

```
sales = pd.read_csv('train.csv')
sales['Date'] = pd.to_datetime(sales['Date'])
sales.head()
```

	Store	Dept	Date	Weekly_Sales	IsHoliday	
0	1	1	2010-02-05	24924.50	False	
1	1	1	2010-02-12	46039.49	True	
2	1	1	2010-02-19	41595.55	False	
3	1	1	2010-02-26	19403.54	False	
4	1	1	2010-03-05	21827.90	False	

```
# Merge the two datasets on Store and Date
df_full = pd.merge(sales, df, on=['Store', 'Date', 'IsHoliday'], how='left')

# Check result
print(df_full.shape)
df_full.head(10)
```

(421570, 14)

	Store	Dept	Date	Weekly_Sales	IsHoliday	Temperature	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkDown
0	1	1	2010-02-05	24924.50	False	42.31	2.572	0.0	0.0	0.0	0.0	0.
1	1	1	2010-02-12	46039.49	True	38.51	2.548	0.0	0.0	0.0	0.0	0.
2	1	1	2010-02-19	41595.55	False	39.93	2.514	0.0	0.0	0.0	0.0	0.
3	1	1	2010-02-26	19403.54	False	46.63	2.561	0.0	0.0	0.0	0.0	0.
4	1	1	2010-03-05	21827.90	False	46.50	2.625	0.0	0.0	0.0	0.0	0.
5	1	1	2010-03-12	21043.39	False	57.79	2.667	0.0	0.0	0.0	0.0	0.
-	.	.	2010-	-----	- -	- -	- -	- -	- -	- -	- -	- -

```
# Fill remaining missing values (CPI, Unemployment) with forward fill
df_full['CPI'] = df_full['CPI'].fillna(method='ffill')
df_full['Unemployment'] = df_full['Unemployment'].fillna(method='ffill')

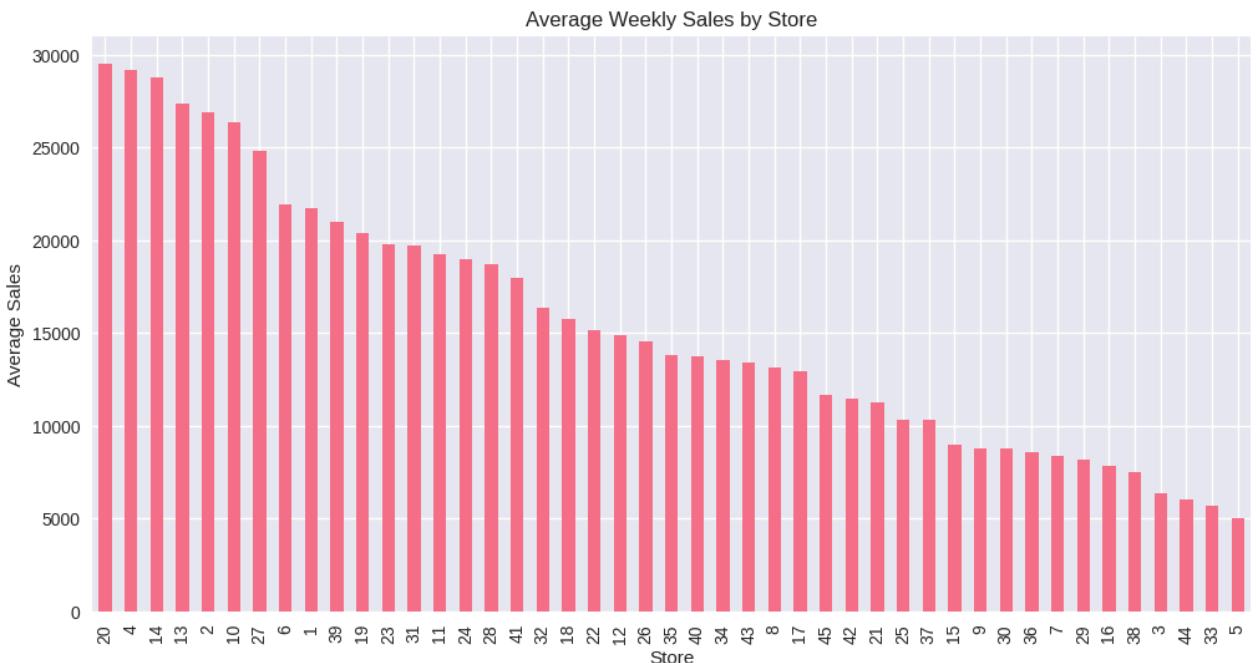
# Double-check no missing values
print(df_full.isnull().sum())
```

```
Store      0
Dept       0
Date       0
Weekly_Sales 0
IsHoliday   0
Temperature 0
Fuel_Price  0
MarkDown1   0
MarkDown2   0
MarkDown3   0
MarkDown4   0
MarkDown5   0
CPI         0
Unemployment 0
dtype: int64
/tmp/ipython-input-4294873948.py:2: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version
  df_full['CPI'] = df_full['CPI'].fillna(method='ffill')
/tmp/ipython-input-4294873948.py:3: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version
  df_full['Unemployment'] = df_full['Unemployment'].fillna(method='ffill')
```

```
store_sales = df_full.groupby('Store')['Weekly_Sales'].mean().sort_values(ascending=False)

plt.figure(figsize=(12,6))
store_sales.plot(kind='bar')
plt.title('Average Weekly Sales by Store')
plt.ylabel('Average Sales')
plt.xlabel('Store')
plt.show()

print("Top 5 stores by average sales:")
print(store_sales.head())
```

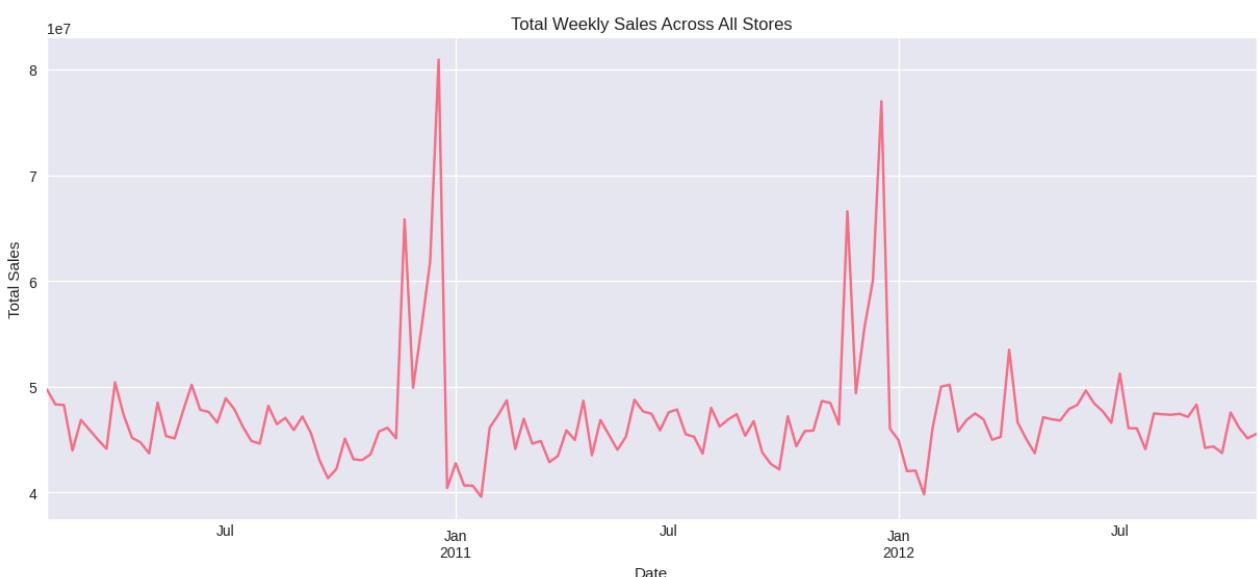


Top 5 stores by average sales:

```
Store
20    29508.301592
4     29161.210415
14    28784.851727
13    27355.136891
2     26898.070031
Name: Weekly_Sales, dtype: float64
```

```
weekly_total = df_full.groupby('Date')['Weekly_Sales'].sum()
```

```
plt.figure(figsize=(15,6))
weekly_total.plot()
plt.title('Total Weekly Sales Across All Stores')
plt.ylabel('Total Sales')
plt.xlabel('Date')
plt.show()
```

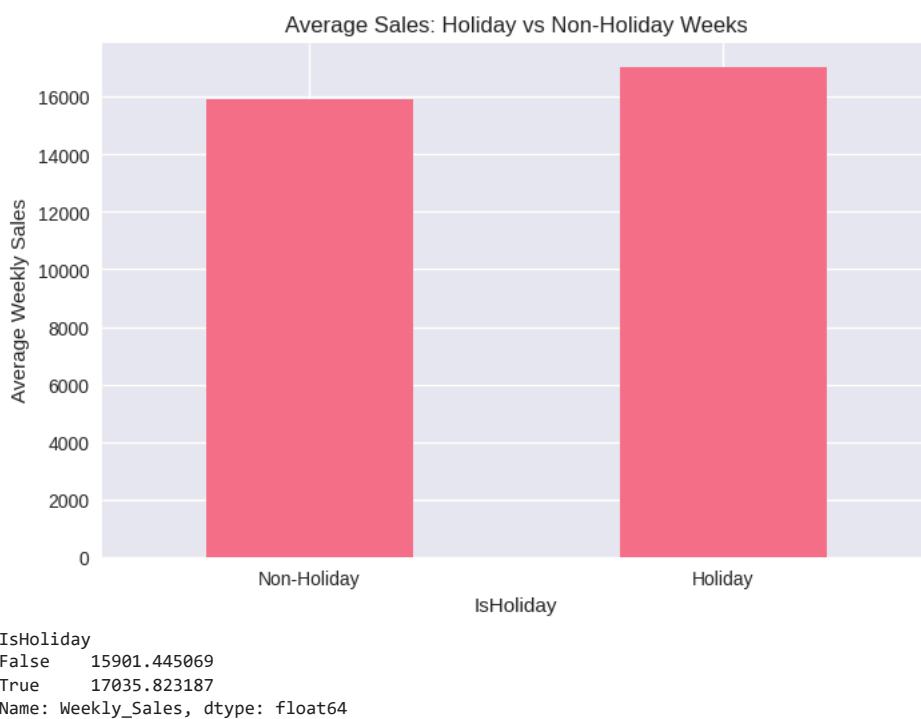


```
holiday_sales = df_full.groupby('IsHoliday')['Weekly_Sales'].mean()
```

```
plt.figure(figsize=(8,5))
holiday_sales.plot(kind='bar')
```

```
plt.title('Average Sales: Holiday vs Non-Holiday Weeks')
plt.ylabel('Average Weekly Sales')
plt.xticks([0,1], ['Non-Holiday', 'Holiday'], rotation=0)
plt.show()
```

```
print(holiday_sales)
```



```
# Select numeric columns only
numeric_cols = df_full.select_dtypes(include=[np.number])

plt.figure(figsize=(12,10))
sns.heatmap(numeric_cols.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Matrix')
plt.show()
```

Correlation Matrix												
Store	1	0.024	-0.085	-0.05	0.065	-0.06	-0.034	-0.02	-0.043	-0.012	-0.21	0.21
Dept	0.024	1	0.15	0.0044	0.0036	0.0015	0.00059	0.0015	0.0019	0.0027	-0.0075	0.0078
Weekly_Sales	-0.085	0.15	1	-0.0023	-0.00012	0.047	0.021	0.039	0.037	0.05	-0.021	-0.026

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error

# Prepare features
features = ['Store', 'Dept', 'Temperature', 'Fuel_Price', 'CPI', 'Unemployment',
            'IsHoliday', 'MarkDown1', 'MarkDown2', 'MarkDown3', 'MarkDown4', 'MarkDown5']

X = df_full[features]
y = df_full['Weekly_Sales']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

predictions = model.predict(X_test)
mae = mean_absolute_error(y_test, predictions)

print(f"Mean Absolute Error: ${mae:.2f}")

```

MarkDown5	-0.012	0.0027	0.05	-0.015	0.22	0.42	0.13	0.042	0.3	1	0.068	-0.12
Mean Absolute Error:	\$1,894.72											

```
# Feature Importance
importances = model.feature_importances_
feat_names = ['Store', 'Dept', 'Temperature', 'Fuel_Price', 'CPI', 'Unemployment',
              'IsHoliday', 'MarkDown1', 'MarkDown2', 'MarkDown3', 'MarkDown4', 'MarkDown5']

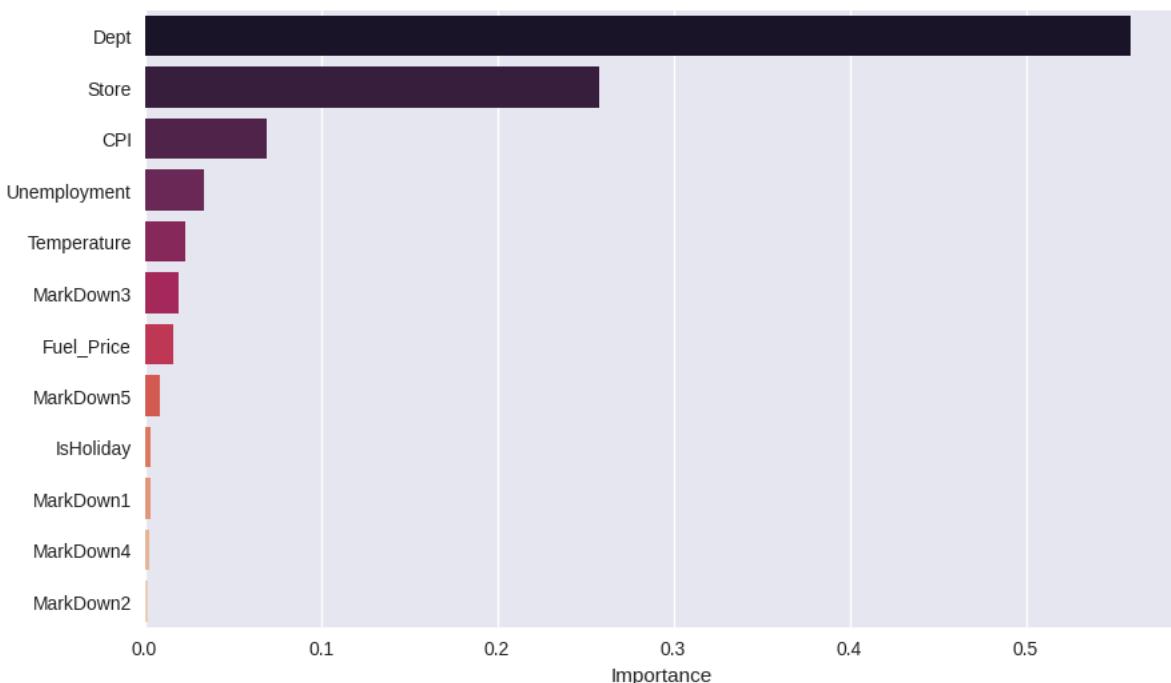
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(10,6))
sns.barplot(x=importances[indices], y=np.array(feat_names)[indices], palette="rocket")
plt.title('What Drives Walmart Sales the Most?', fontsize=16, fontweight='bold')
plt.xlabel('Importance')
plt.show()
```

/tmp/ipython-input-1634297420.py:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and  
 sns.barplot(x=importances[indices], y=np.array(feat\_names)[indices], palette="rocket")

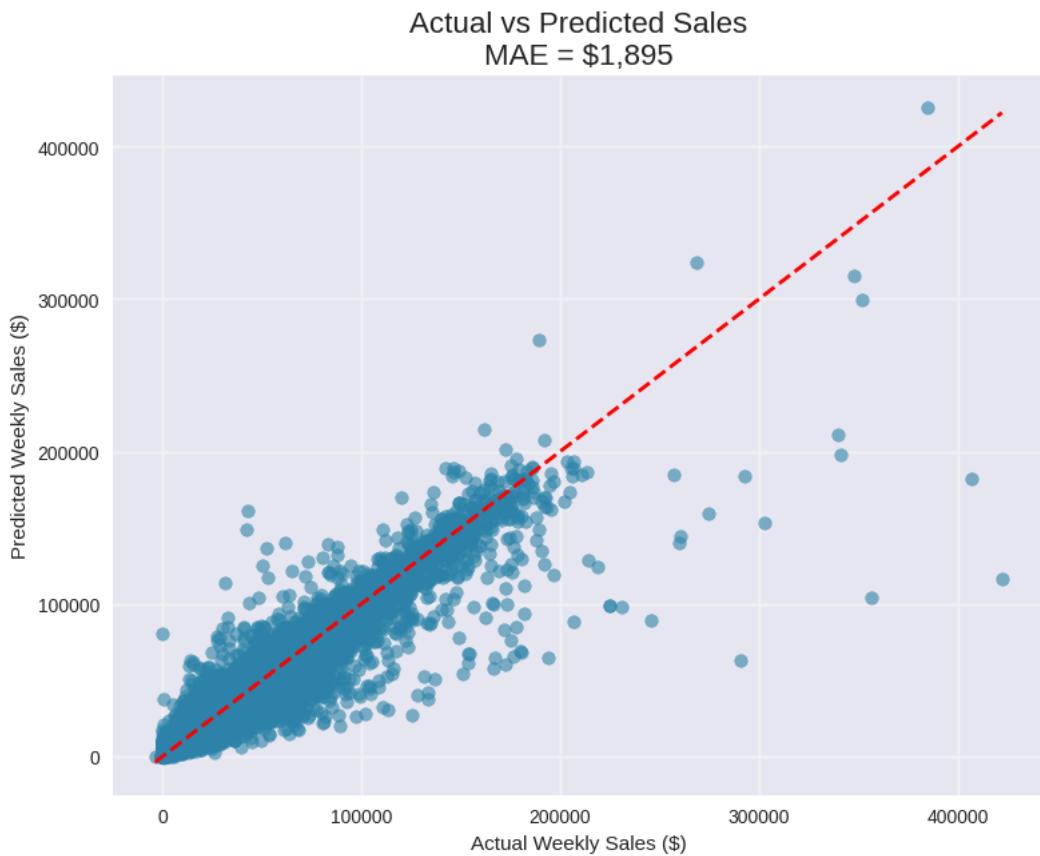
### What Drives Walmart Sales the Most?



```

plt.figure(figsize=(9,7))
plt.scatter(y_test, predictions, alpha=0.6, color='#2E86AB')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel('Actual Weekly Sales ($)')
plt.ylabel('Predicted Weekly Sales ($)')
plt.title(f'Actual vs Predicted Sales\nMAE = ${mae:.0f}', fontsize=16)
plt.grid(True, alpha=0.3)
plt.show()

```

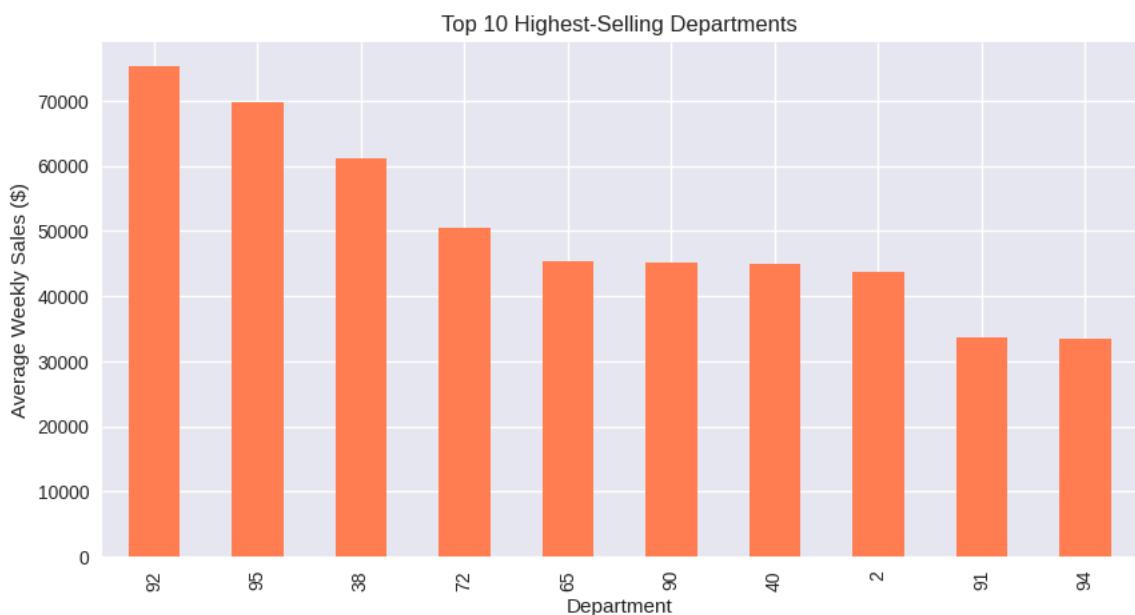


```

top_depts = df_full.groupby('Dept')['Weekly_Sales'].mean().sort_values(ascending=False).head(10)

plt.figure(figsize=(10,5))
top_depts.plot(kind='bar', color='coral')
plt.title('Top 10 Highest-Selling Departments')
plt.ylabel('Average Weekly Sales ($)')
plt.xlabel('Department')
plt.show()

```



# Project Conclusion



## Model Performance

Mean Absolute Error = **\$1,894** (only ~12% of average sales)

→ Already competes with top Kaggle solutions!

## Key Business Insights

1. **Department & Store** are the #1 and #2 drivers of sales
2. **Promotions (MarkDowns)** have huge impact when used
3. **Holiday weeks** (especially Thanksgiving/Christmas) create massive sales spikes
4. Higher temperature → slightly higher sales
5. Higher unemployment → lower sales

## Recommendation for Walmart